

Product Data Explorer — Full-Stack Assignment

Goal

Build a production-minded product exploration platform that lets users navigate from high-level headings → categories → products → product detail pages powered by **live, on-demand scraping**.

All scraping **must be done against [World of Books](#)**.

- ★ You need to use the technologies mentioned in this document.

Candidates must submit:

- **GitHub repo link** (public or private with access)
 - **Deployed project link** (live frontend + working backend)
 - **Both links should be submitted through the following Google Form:**
<https://forms.gle/AiZRVZL2tyoQSups5> .
-

High-level requirements (must have)

Frontend

- Tech: **React (Next.js, App Router), TypeScript, Tailwind CSS**.
- Core pages/components:
 - Landing / Home (shows navigation headings)
 - Category drilldown pages
 - Product grid / results (supports paging / limit)
 - Product detail page (reviews, ratings, recommendations, metadata)
 - About / Contact / README page in site
- UX:

- Responsive (desktop & mobile), accessible (WCAG AA basics), skeleton/loading states and smooth transitions.
- Persist user navigation & browsing history client-side and via backend for reloading.
- Use a client data fetching strategy (SWR or React Query recommended).
- Deliverables:
 - Live URL to the deployed frontend
 - Instructions in README: how to run locally, environment variables, build steps

Backend

- Tech: **NestJS (Node + TypeScript)**.
- Production-ready DB: PostgreSQL, MongoDB, or other (must justify choice).
- Expose RESTful endpoints (see example API contract below).
- On relevant calls, trigger a **real-time scrape** (Crawlee + Playwright) and store results. Allow:
 - On-demand scrapes triggered by user actions
 - Safe caching to avoid excessive scraping
- Robust engineering:
 - Proper DTO validation, error handling, logging, and resource cleanup
 - Concurrency handling, deduplication of scrape results, idempotency where applicable
 - Rate limiting / backoff for external sites and queueing of long running scrapes

Scraping (World of Books)

- Target site: <https://www.worldofbooks.com/>
- Use **Crawlee + Playwright** (or equivalent headless browser framework).

- Extract and persist:
 - Navigation headings (e.g., “Books”, “Categories”, “Children’s Books”, etc.)
 - Categories & subcategories
 - Product tiles/cards:
 - Title, Author, Price, Image, Product Link, Source ID
 - Product detail pages:
 - Full description
 - User reviews & ratings (if present)
 - Related/recommended products
 - Any additional available metadata (publisher, publication date, ISBN, etc.)
- Save all scraped data into your DB with relationships & unique constraints.
- Implement **deduplication** and **caching with expiry** so repeated scrapes don’t overload World of Books.
- Provide a way to re-fetch updated product data on demand.

Ethical scraping reminder:

- Respect robots.txt and terms of service.
- Use proper rate limiting and delays.
- Implement retries and exponential backoff.
- Cache results wherever possible to avoid hitting the site repeatedly.

Suggested Database schema (entities)

- `navigation` — id, title, slug, last_scraped_at

- `category` — id, navigation_id, parent_id, title, slug, product_count, last_scraped_at
 - `product` — id, source_id, title, price, currency, image_url, source_url, last_scraped_at
 - `product_detail` — product_id (FK), description, specs (json), ratings_avg, reviews_count
 - `review` — id, product_id, author, rating, text, created_at
 - `scrape_job` — id, target_url, target_type, status, started_at, finished_at, error_log
 - `view_history` — id, user_id (optional), session_id, path_json, created_at
 - Add indexes on `source_id`, `last_scraped_at`, and unique constraints on `source_url/source_id`.
-

Non-functional requirements

- **Security:** sanitize inputs, secure environment variables, do not commit secrets, enable CORS properly, minimal rate limiting.
 - **Performance & caching:** caching layer in DB (or Redis) with explicit expiry; avoid re-scraping unchanged pages.
 - **Observability:** logging, basic metrics, and error tracking (console + file or a service).
 - **Reliability:** queue/worker model for scrapes (do not block request thread); idempotent jobs.
 - **Accessibility:** semantic HTML, keyboard nav, alt on images, color contrast.
-

Deliverables

1. **GitHub repo link** with:
 - `frontend/` and `backend/` folders
 - CI pipeline (GitHub Actions) for lint/test/build (recommended)

- README with architecture overview, design decisions, and deployment instructions
- Database schema and sample seed script
- API documentation (Swagger or markdown)
- Tests (unit + a couple integration tests)
- Dockerfiles (bonus, but preferred)

2. Deployed project link(s):

- Frontend URL (production)
 - Note: Must be live at submission time
-

Acceptance checklist (must pass)

- Landing loads navigation headings (from World of Books via backend)
 - Drilldown loads categories/subcategories (from World of Books via backend)
 - Product grid displays real products (scraped from World of Books)
 - Product detail page includes description, reviews/ratings, recommendations (scraped from World of Books)
 - DB persists all scraped objects reliably
 - On-demand scrape can refresh a product/category
 - Frontend responsive and accessible baseline
 - README + deploy links + API docs present
 - Repo builds and runs with the provided instructions
-

Evaluation rubric (weights)

- **Correctness & completeness (35%)** — feature coverage vs requirements
 - **Architecture & engineering quality (20%)** — code structure, DTOs, validation, error handling
 - **Scraping reliability & design (15%)** — safe scraping, queueing, dedupe, caching
 - **UX & accessibility (10%)** — responsiveness, loading states, accessibility basics
 - **Docs & deploy (10%)** — README, API docs, deployed links working
 - **Tests & CI (10%)** — basic tests and CI pipeline
-

Bonus (highly valued)

- Product search + rich filters (price range, rating, author)
 - Intelligent caching / refresh strategy (DB-backed TTL, conditional scraping)
 - SWR / React Query with optimistic UI updates
 - Personalized recommendations or similarity engine (simple content-based OK)
 - Full Docker setup for frontend, backend, DB + `docker-compose`
 - Comprehensive test coverage (unit + e2e)
 - API versioning and OpenAPI/Swagger with examples
 - CI-based deploy to Vercel / Render / Heroku / Railway / Fly.io
-

Tips & constraints for candidates

- **Be kind to World of Books:** implement delays, backoff, and cache results.
- **Focus on core features first** before optional bonuses.
- **Do not commit secrets** — use `.env.example`.

- **Include a fallback seed script** so reviewers can test even if scraping fails during review.
-